

MITCHELL WILLIAMS

AI SOLUTIONS ARCHITECTURE · GOVERNED-VOICE AGENT SYSTEMS · MULTI-AGENT ORCHESTRATION

Seattle, WA (open to relocation) | 812-219-6077 | mitwilli@gmail.com
linkedin.com/in/mitwilli | github.com/mitwilli-create | thestorytellermitch.com

SUMMARY

Solutions architect who builds governed-voice AI systems: agents whose output is measured, gated, and citation-bound by architecture. Governed voice is a pattern I have shipped twice: **voice-os** scores writing on six measured axes against a codified voice "Constitution" with a banned-phrase checker, and the executive **digital twin** I shipped at Google drafts VP communications in a governed voice, cutting review back-and-forth by about 70% (internal, self-measured).

Eight years at Google, the last two running the full internal solutions lifecycle for an internal customer community of 1,000+ Principal, Distinguished, and Fellow engineers. At Google Cross-Google Engineering (xGE) I ran discovery, build, deployment, and handoff: two production agents shipped solo, VP-level adoption, a full **operator runbook**. Internal customers, not external accounts.

Before Google, a decade of live national news, where requirements arrive mid-broadcast and the system holds or it does not. The same discipline governs my agent design: hard guarantees over vibes, refusal over guessing.

PROJECTS

Repos public at github.com/mitwilli-create.

voice-os · Six-axis measured voice scoring, dual-persona routing, and QA gates calibrated on a multi-million-word corpus; codified voice "Constitution" plus banned-phrase checker. The personal build of the governed-voice technique behind the Google executive digital twin.

tax-verification-agent · Citation-gated verification over a four-layer IRS knowledge base on Claude: no claim without a source, unknowns surface as placeholders instead of guesses. Caught a ~\$19K multi-state error that filing software pre-filled and defended as correct.

monolith · TypeScript agent built on hard guarantees: a deterministic constraint-gate contract that returns INSUFFICIENT_DATA instead of guessing when the inputs cannot support an answer.

career-ops · Multi-agent system running my job search daily: multi-LLM council-of-models with adjudication, custom Claude Code skills and MCP servers, 52 scheduled agents, 1,150+ scored role reports, live dashboard under a unified spend cap. CI-green; public on GitHub.

PictureLock · Seven-stage AI video pipeline in Node on the ElevenLabs stack; live demo receipt: a 53-second creative-directed short; the committed run manifest carries \$8.26 for the artifacts in the cut, the Spanish dub a separately logged \$0.42, every call logged.

FIT FOR AI SOLUTIONS ARCHITECTURE

- **Voice fidelity as a testable primitive.** The executive digital twin measures every draft against a voice "Constitution" and banned-phrase checklist; it cut VP review back-and-forth by about 70% (internal, self-measured).
- **Hard guarantees over vibes.** The tax-verification agent is citation-gated by architecture and caught a ~\$19K error that commercial software defended as correct. Monolith returns INSUFFICIENT_DATA instead of guessing.

EXPERIENCE

Internal Communications Lead

Google · Cross-Google Engineering (xGE) | June 2024 – Present

Communications lead and program manager for xGE, the horizontal org that drives technical change across Google.

Communications-triage agent

- **Problem:** High-volume comms intake for the senior-IC community consumed operator hours on low-complexity requests.
- **Action:** Built and deployed solo (Apps Script + Gemini) a three-prompt pipeline (triage, revise, escalate) with conditional knowledge-base loading and tiered routing; human-in-the-loop authority holds on every High-touch classification.
- **Impact:** Documented end-to-end and handed off with a full operator runbook a non-builder can run and extend.



Executive-communications digital twin

- **Problem:** High-stakes VP communications required long review cycles to hold a consistent voice.
- **Action:** Built a Gemini agent loaded with a "Constitution" of voice rules, a banned-phrase verification checklist, and engagement-summary generators.
- **Impact:** Review cycles with the VP dropped by about 70%.

MEJ (My Engineering Journey), AI mentorship-matching agent

- **Problem:** Manual mentorship matching across the senior-engineer community capped program capacity.
- **Action:** Architected and prototyped the matching agent and app as lead PM and architect: Python matching engine, Gemini prompt-chaining pipelines, agent security guardrails.
- **Impact:** Working prototype, framed as exactly that; production numbers wait until the system earns them.

Approvals governance overhaul

- **Problem:** Cross-org technical mandates stalled without a clear approvals framework.
- **Action:** Authored the Approvals User Guide (v3, "Effectively Approved"), the Approvals Matrix, and the Roles & Responsibilities framework.
- **Impact:** Drove named cross-org mandates (Keystore Rotation, Memory Safety TRR) to approval.

Senior Communications & Content Strategist

Google · Corporate Engineering (TechStop) | April 2018 – June 2024

Owned communications and content enabling technical outcomes for a global, mixed-technical audience across Google's internal IT org and TechStop, its global help desk; directed the Q1-2020 global remote-work shift strategy, driving an 80% increase in self-provisioning efficiency the following quarter.

2023 self-help content audit

- **Problem:** Ticket volume included issues users could have solved themselves.
- **Action:** Audited the self-help corpus and rewrote the highest-traffic, lowest-solve articles.
- **Impact:** Self-service solvability up 9.03%, roughly 10,670 fewer tickets per year.

EARLIER CAREER: LIVE NATIONAL NEWS (2007 – 2018)

The engine of live news is the engine of solutions architecture: messy human requirements translated, live and under deadline, into boundaries a system can hold.

- **AJ+, Senior Producer (2016 – 2018):** Led 10+ producers; designed a day-of explainer production line (proof-of-concept: 50M+ views) that became a permanent format.
- **Fusion (ABC News / Univision), Senior Producer, "America With Jorge Ramos" (2013 – 2015):** Owned daily line-production on the primetime flagship through a 179% viewership-growth window; field-produced live from Hong Kong's [Umbrella Revolution](#).
- **HuffPost Live, Segment Producer (2012 – 2013):** Produced through the award-winning launch era (Mashable "Biggest Innovation in Media," 2012; Webby, 2013); sustained coverage under active litigation posture, no retractions.
- **Al Jazeera English "The Stream" · CGTN America (2010 – 2012):** Founding-team Associate Producer; launch night (May 2, 2011; 250M households) pivoted to the [bin Laden](#) raid and pioneered social-data journalism.

EDUCATION

BA, Journalism | Indiana University Bloomington

CERTIFICATIONS

AI Fluency · Claude 101 · Agent Skills · Model Context Protocol · Skill Builder (Opus cohort) | Anthropic | 2026

SKILLS

AI / Agent Engineering & Build: multi-provider LLM orchestration (Anthropic, OpenAI, Google, xAI, Perplexity), custom MCP servers, Claude Code skills, eval and regression gates, RAG + vector stores, prompt engineering, LangGraph / LangChain, Node.js, Playwright, launchd fleet operations, Apps Script, Gemini API, Python.

Program & Communications: governance and approvals frameworks, intake protocols, metrics design, stakeholder management, executive communications at engineering scale.